

Business Analytics

Decision Tree Classifiers

Prof. Bichler

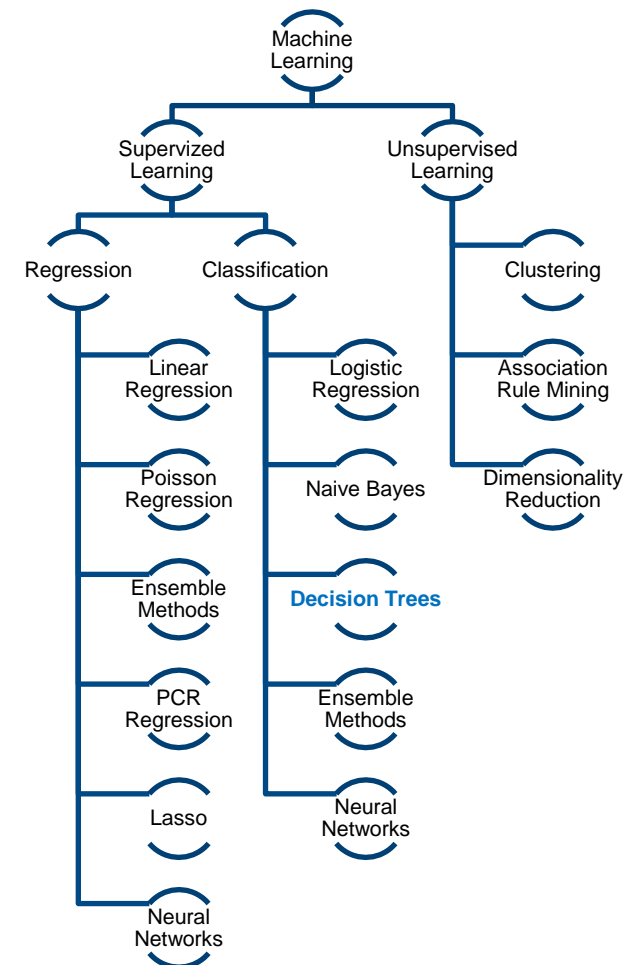
Decision Sciences & Systems

Department of Informatics

Technische Universität München


Course Content

- Introduction
- Regression Analysis
- Regression Diagnostics
- Logistic and Poisson Regression
- Naive Bayes and Bayesian Networks
- **Decision Tree Classifiers**
- Data Preparation and Causal Inference
- Model Selection and Learning Theory
- Ensemble Methods and Clustering
- High-Dimensional Problems
- Association Rules and Recommenders
- Neural Networks



Recommended Literature

- **Data Mining: Practical Machine Learning Tools and Techniques**
 - Ian H. Witten, Eibe Frank, Mark A. Hall, Christopher Pal.
 - <http://www.cs.waikato.ac.nz/ml/weka/book.html>
 - Section: 4.3, 6.1
- **Alternative literature**
 - Machine Learning
 - Tom M. Mitchell, 1997
 - Data Mining: Introductory and Advanced Topics
 - Margaret H. Dunham, 2003



Outline for Today

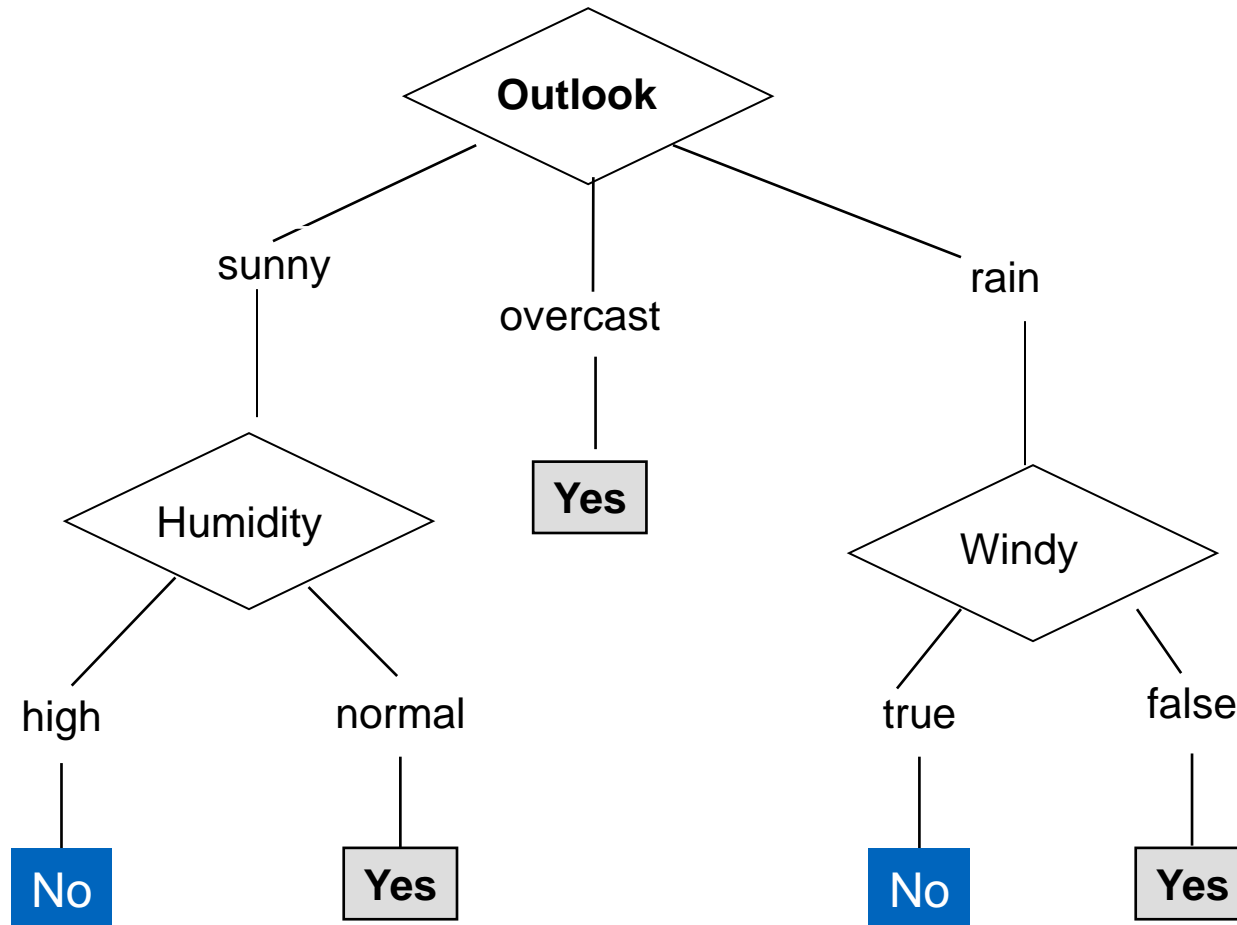
- **Choosing a splitting attribute in decision trees**
 - Information gain
 - Gain ratio
 - Gini index
- Numeric attributes
- Missing values
- Relational rules



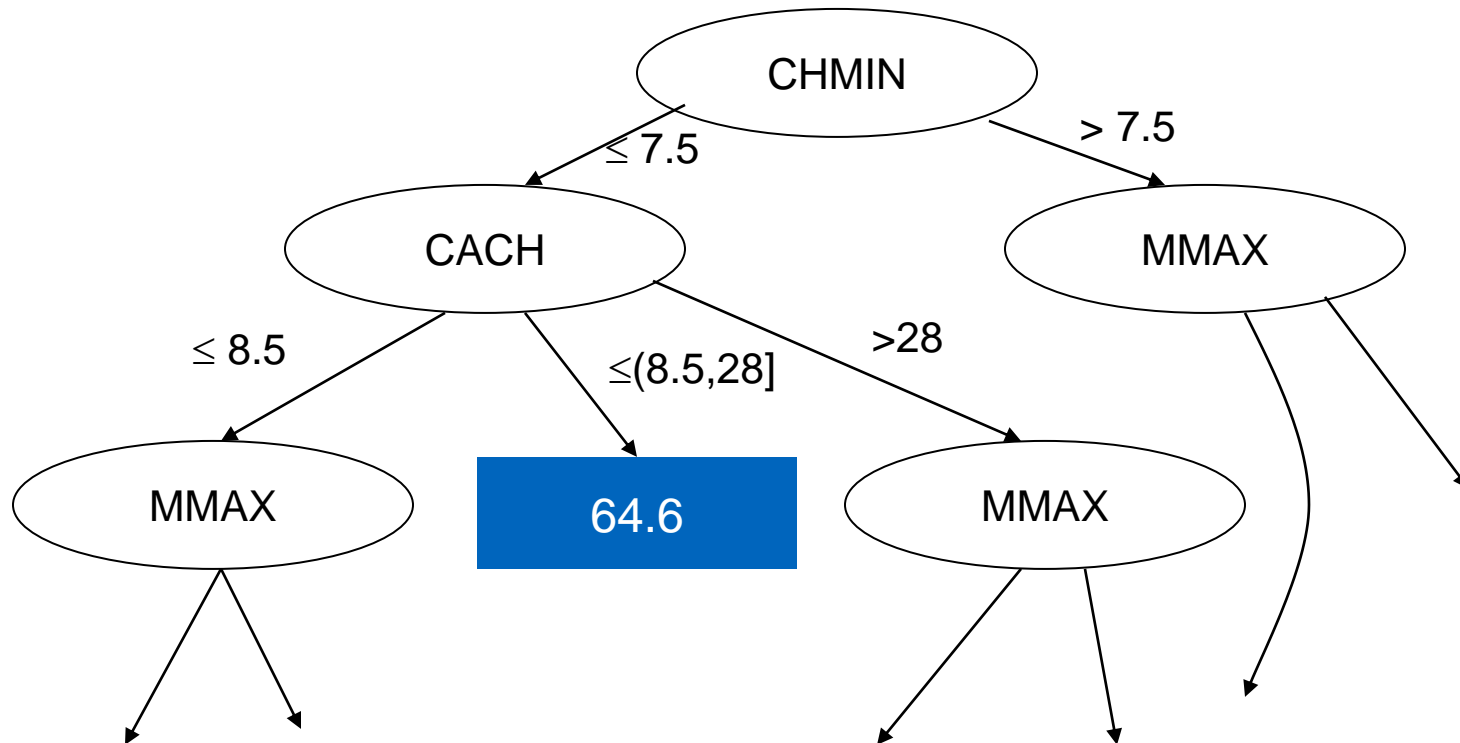
The Weather Data

Outlook	Temp	Humidity	Windy	Play
Sunny	Hot	High	False	No
Sunny	Hot	High	True	No
Overcast	Hot	High	False	Yes
Rainy	Mild	High	False	Yes
Rainy	Cool	Normal	False	Yes
Rainy	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Sunny	Mild	High	False	No
Sunny	Cool	Normal	False	Yes
Rainy	Mild	Normal	False	Yes
Sunny	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Rainy	Mild	High	True	No

Example Tree for “Play?”

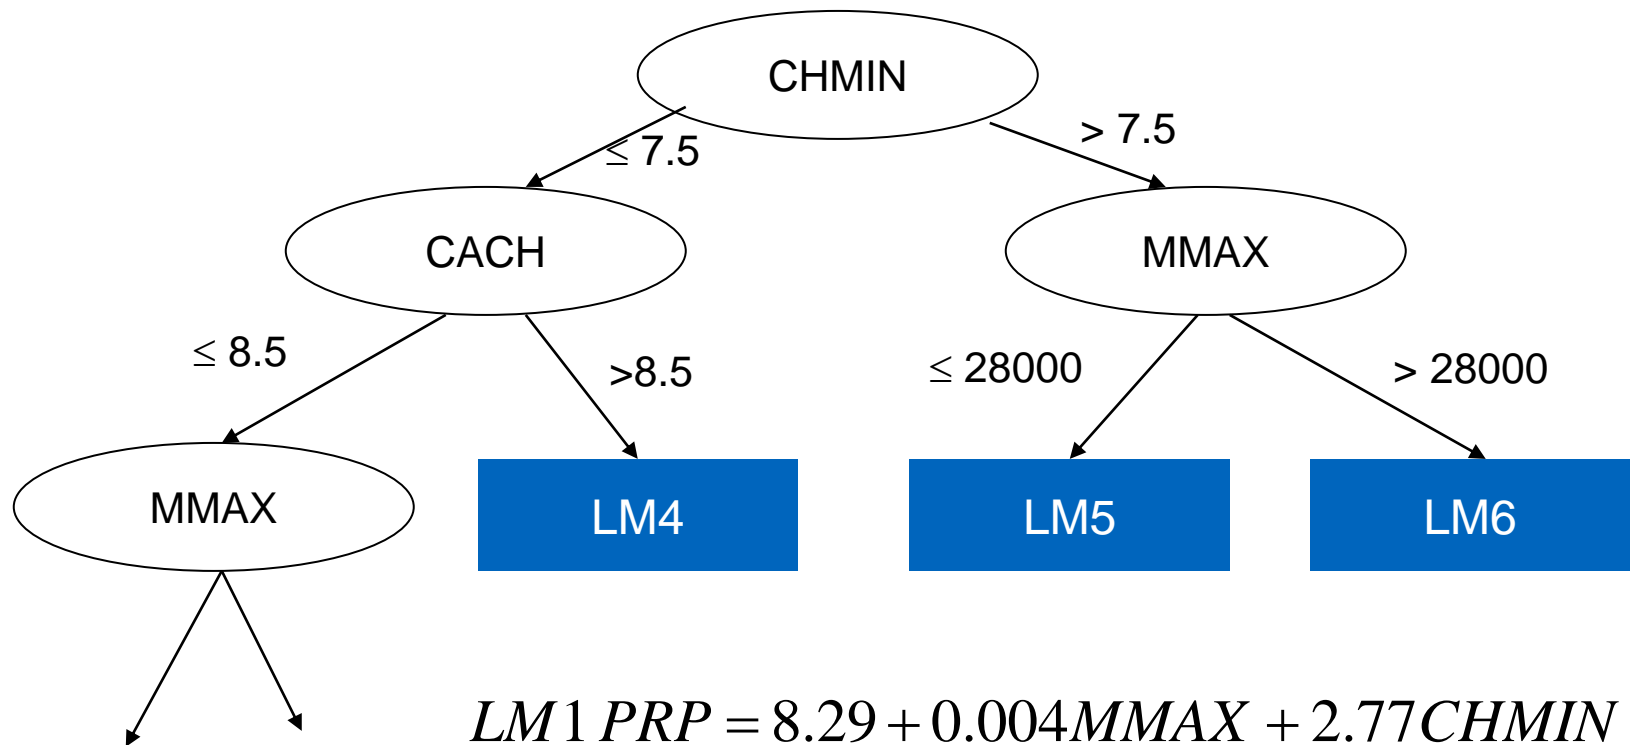


Regression Tree



- High Accuracy
- Large and possibly awkward

Model Trees



$$LM1 PRP = 8.29 + 0.004MMAX + 2.77CHMIN$$

$$LM2 PRP = \dots$$

⋮

Decision Trees

- An internal node is a test on an attribute
- A branch represents an outcome of the test, e.g., *Color = red*
- A leaf node represents a class label or class label distribution
- At each node, one attribute is chosen to split training examples into distinct classes as much as possible
- A new case is classified by following a matching path to a leaf node

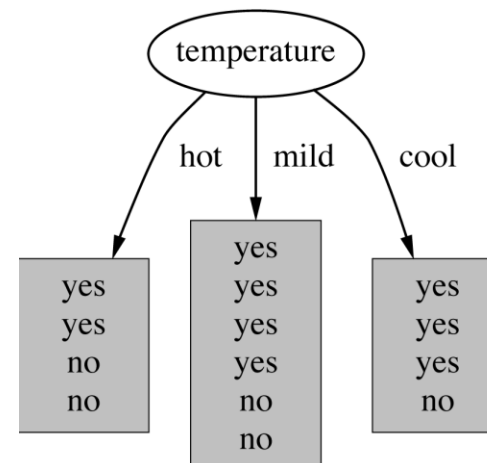
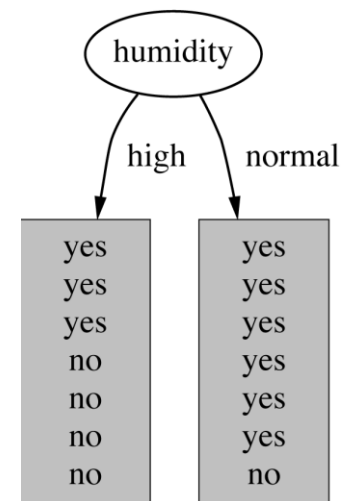
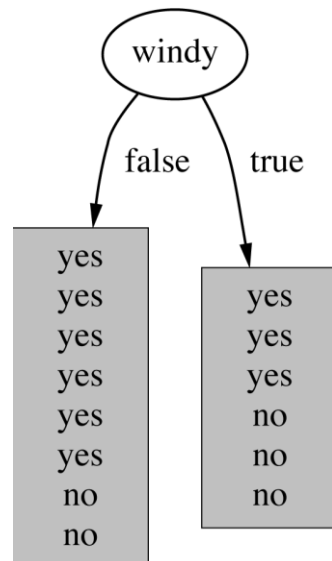

Building Decision Tree

- Top-down tree construction
 - At start, all training examples are at the root
 - Partition the examples recursively by choosing one attribute each time
- Bottom-up tree pruning
 - Remove subtrees or branches, in a bottom-up manner, to improve the estimated accuracy on new cases

Choosing the Splitting Attribute

- At each node, available attributes are evaluated on the basis of separating the classes of the training examples
- A “goodness” function is used for this purpose
- Typical goodness functions:
 - information gain (ID3/C4.5)
 - information gain ratio
 - gini index (CART)

Which Attribute to Select?



A Criterion for Attribute Selection


- Which is the best attribute?
 - The one which will result in the smallest tree
 - Heuristic: choose the attribute that produces the “purest” nodes
- Popular impurity criterion: information gain
 - Information gain increases with the average purity of the subsets that an attribute produces
- Strategy: choose attribute that results in greatest information gain

Computing Information

- Information is measured in *bits*
 - Given a probability distribution, the info required to predict an event, i.e. if play is yes or no, is the distribution's *entropy*
 - *Entropy* gives the information required in bits (this can involve fractions of bits!)
- Formula for computing the information entropy:

$$\text{entropy}(p_1, p_2, \dots, p_n) = -p_1 \log_2 p_1 - p_2 \log_2 p_2 \dots - p_n \log_2 p_n$$

Example: attribute “Outlook”



“Outlook” = “Sunny”:

$$\text{info}([2,3]) = \text{entropy}(2/5, 3/5) = -2/5 \log_2(2/5) - 3/5 \log_2(3/5) = 0.971 \text{ bits}$$

“Outlook” = “Overcast”:

$$\text{info}([4,0]) = \text{entropy}(1,0) = -1 \log_2(1) - 0 \log_2(0) = 0 \text{ bits}$$



Note: $\log(0)$ is not defined, but we evaluate $0 \cdot \log(0)$ as zero

“Outlook” = “Rainy”:

$$\text{info}([3,2]) = \text{entropy}(3/5, 2/5) = -3/5 \log_2(3/5) - 2/5 \log_2(2/5) = 0.971 \text{ bits}$$

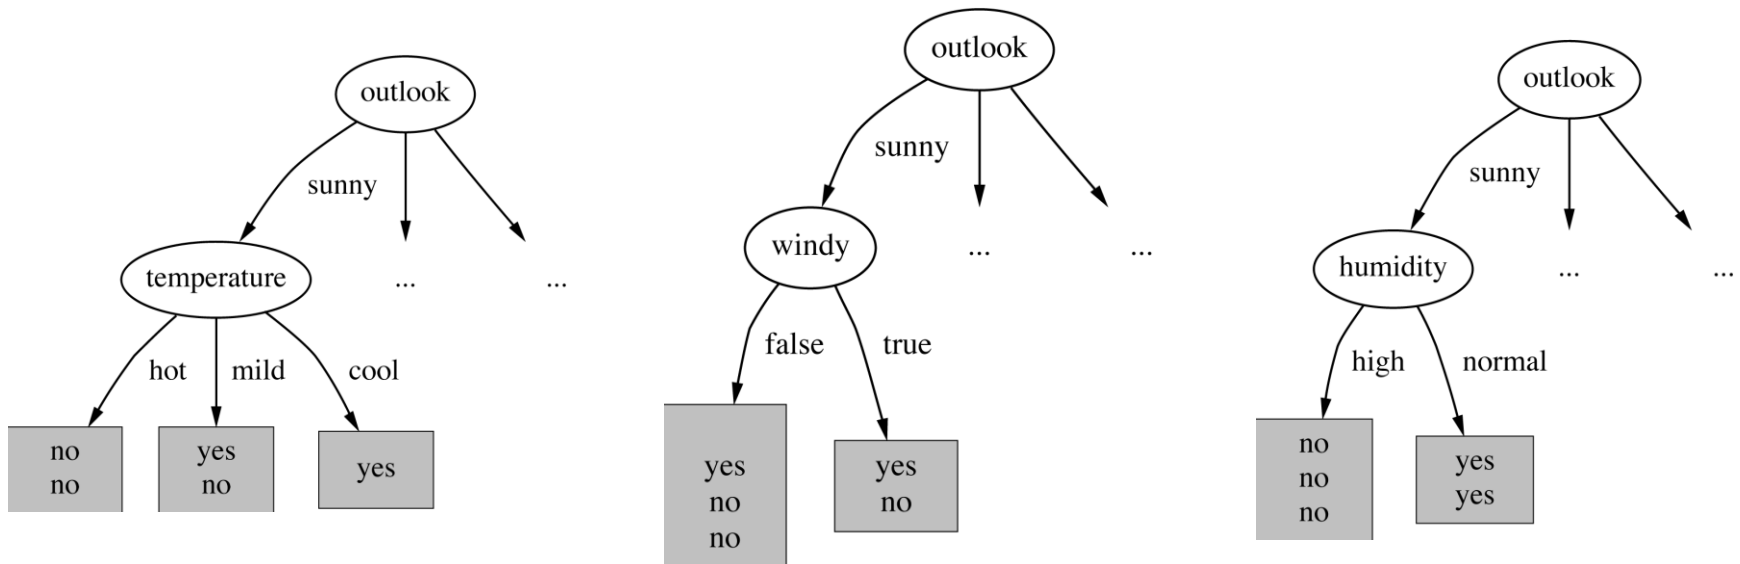
Expected information for attribute:

$$\begin{aligned} \text{info}([2,3], [4,0], [3,2]) &= (5/14) \times 0.971 + (4/14) \times 0 + (5/14) \times 0.971 \\ &= 0.693 \text{ bits} \end{aligned}$$

Computing the Information Gain

- Information gain: (information before split) – (information after split)
 $\text{gain}(\text{Outlook}) = \text{info}([9,5]) - \text{info}([2,3], [4,0], [3,2])$
 $= 0.940 - 0.693 = 0.247 \text{ bits}$
- Information gain for attributes from weather data:
 $\text{gain}(\text{Outlook}) = 0.247 \text{ bits}$
 $\text{gain}(\text{Temperature}) = 0.029 \text{ bits}$
 $\text{gain}(\text{Humidity}) = 0.152 \text{ bits}$
 $\text{gain}(\text{Windy}) = 0.048 \text{ bits}$

Continuing to Split

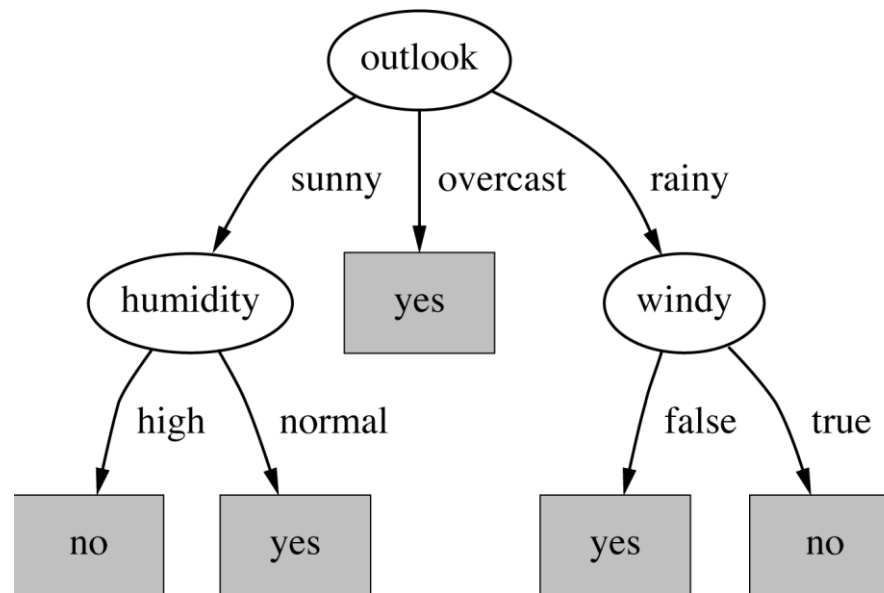


$\text{gain}(\text{Temperatur}) = 0.571 \text{ bits}$

$\text{gain}(\text{Humidity}) = 0.971 \text{ bits}$

$\text{gain}(\text{Windy}) = 0.020 \text{ bits}$

The Final Decision Tree



- Note: not all leaves need to be pure;
⇒ Splitting stops when data can't be split any further

Wish List for a Purity Measure

- Properties we require from a purity measure:
 - When node is pure, measure should be zero ($=0$)
 - When impurity is maximal (i.e. all classes equally likely), measure should be maximal (e.g., 1 for boolean values)
 - Multistage property: $\text{info}[2,3,4] = \text{info}[2,7] + 7/9 \text{ info}[3,4]$
- Entropy is a function that satisfies these properties

The diagram shows the entropy formula $entropy(S) = \sum_{i=1}^n -p_i \log_2 p_i$ with three annotations. An arrow points from a box labeled 'Training data (instances)' to the S in the function. Another arrow points from a box labeled 'Number of classes' to the n in the summation index. A third arrow points from a box labeled 'Probability of S being classified to i ' to the p_i in the formula. To the right of the formula, a note in parentheses states '(scales from 0 to $\max \log_2 n$)'.

$$entropy(S) = \sum_{i=1}^n -p_i \log_2 p_i \quad (\text{scales from 0 to } \max \log_2 n)$$

Annotations:

- Training data (instances) points to S .
- Number of classes points to n .
- Probability of S being classified to i points to p_i .

Entropy

- Entropy in general, describes the randomness of a system
 - entropy = 0 describes a perfectly ordered system
- The term was used in thermodynamics and statistical mechanics
- „A Mathematical Theory of Communication“ by Claude Shannon 1948 defines entropy and information theory
 - „uncertainty about the contents of a message“

Claude Shannon

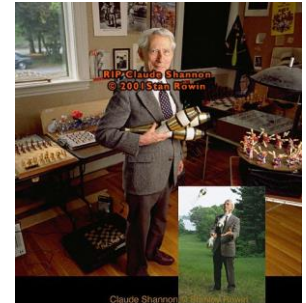
Born: 30 April 1916, Died: 23 February 2001

Shannon is famous for having founded information theory with one landmark paper published in 1948 (*A Mathematical Theory of Communication*).

Information theory was developed to find fundamental limits on compressing and reliably communicating data. Communications over a channel was the primary motivation. Channels (such as a phone line) often fail to produce exact reconstruction of a signal; noise, periods of silence, and other forms of signal corruption often degrade quality. How much information can one hope to communicate over a noisy (or otherwise imperfect) channel?

An important application of information theory is coding theory:

- Data compression (by removing redundancy in data)
- Error-correcting codes add just the right kind of redundancy (i.e. error correction) needed to transmit the data efficiently and faithfully across a noisy channel.



Expected Information Gain

$$\text{gain}(S, a) = \text{entropy}(S) - \sum_{v \in \text{Values}(a)} \frac{|S_v|}{|S|} \text{entropy}(S_v)$$

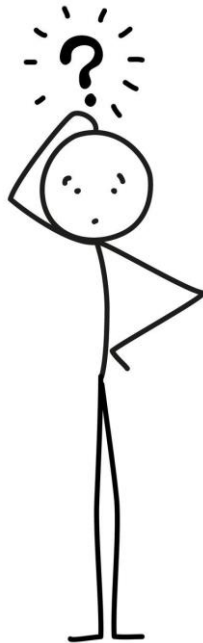
$$S_v = \{s \in S : a(s) = v\}$$

← All possible values
for attribute a

$\text{gain}(S, a)$ is the information gained adding a sub-tree
(Reduction in number of bits needed to classify an instance)

Problems?

Which splitting attribute would be selected in the following example?




ID	Outlook	Temperature	Humidity	Windy	Play?
A	sunny	hot	high	false	No
B	sunny	hot	high	true	No
C	overcast	hot	high	false	Yes
D	rain	mild	high	false	Yes
E	rain	cool	normal	false	Yes
F	rain	cool	normal	true	No
G	overcast	cool	normal	true	Yes
H	sunny	mild	high	false	No
I	sunny	cool	normal	false	Yes
J	rain	mild	normal	false	Yes
K	sunny	mild	normal	true	Yes
L	overcast	mild	high	true	Yes
M	overcast	hot	normal	false	Yes
N	rain	mild	high	true	No

Weather Data with ID Code

ID	Outlook	Temperature	Humidity	Windy	Play?
A	sunny	hot	high	false	No
B	sunny	hot	high	true	No
C	overcast	hot	high	false	Yes
D	rain	mild	high	false	Yes
E	rain	cool	normal	false	Yes
F	rain	cool	normal	true	No
G	overcast	cool	normal	true	Yes
H	sunny	mild	high	false	No
I	sunny	cool	normal	false	Yes
J	rain	mild	normal	false	Yes
K	sunny	mild	normal	true	Yes
L	overcast	mild	high	true	Yes
M	overcast	hot	normal	false	Yes
N	rain	mild	high	true	No

Split for ID Code Attribute



Entropy of split = 0 (since each leaf node is “pure”), having only one case.

Information gain is maximal for ID code

Highly-Branching Attributes

- Problematic: attributes with a large number of values (extreme case: ID code)
- Subsets are more likely to be pure if there is a large number of values
 - Information gain is biased towards choosing attributes with a large number of values
 - This may result in overfitting (selection of an attribute that is non-optimal for prediction)

Gain Ratio

- Gain ratio: a modification of the information gain that reduces its bias on high-branch attributes
- Gain ratio takes number and size of branches into account when choosing an attribute
- It corrects the information gain by taking the intrinsic information of a split into account (i.e. how much info do we need to tell which branch an instance belongs to)

Computing the Gain Ratio

Example: intrinsic information for ID code

$$\text{intrinsic_info}(1,1, \dots, 1) = 14 * \left(-\frac{1}{14} * \log_2 \left(\frac{1}{14} \right) \right) = 3.807 \text{ bits}$$

Importance of attribute decreases as intrinsic information grows.

Example of gain ratio: $\text{gainRatio}(S, a) = \frac{\text{gain}(S,a)}{\text{intrinsic_info}(S,a)}$

Example: $\text{gainRatio}(\text{ID_Code}) = \frac{0.940\text{bits}}{3.807\text{bits}} = 0.246$

Gain Ratios for Weather Data

Outlook		Temperature	
Info:	0.693	Info:	0.911
Gain: $0.940 - 0.693$	0.247	Gain: $0.940 - 0.911$	0.029
Split info: $\text{info}([5,4,5])$	1.577	Split info: $\text{info}([4,6,4])$	1.557
Gain ratio: $0.247/1.577$	0.156	Gain ratio: $0.029/1.557$	0.019

Humidity		Windy	
Info:	0.788	Info:	0.892
Gain: $0.940 - 0.788$	0.152	Gain: $0.940 - 0.892$	0.048
Split info: $\text{info}([7,7])$	1.000	Split info: $\text{info}([8,6])$	0.985
Gain ratio: $0.152/1$	0.152	Gain ratio: $0.048/0.985$	0.049

More on the Gain Ratio

- „Outlook” still comes out top
- However:
 - “ID code” has still greater gain ratio (0.246)
 - Standard fix: ad hoc test to prevent splitting on that type of attribute
- Problem with gain ratio: it may overcompensate
 - May choose an attribute just because its intrinsic information is very low
 - Standard fix:
 - First, only consider attributes with greater than average information gain
 - Then, compare them on gain ratio

The Splitting Criterion in CART

- Classification and Regression Tree (CART)
- Developed 1974-1984 by 4 statistics professors
 - Leo Breiman (Berkeley), Jerry Friedman (Stanford), Charles Stone (Berkeley), Richard Olshen (Stanford)
- Gini Index is used as a splitting criterion
- Both C4.5 and CART are robust tools
- No method is always superior – experiment!

Gini Index for 2 Attribute Values

- For example, two classes, *Pos* and *Neg*, and dataset *S* with *p* *Pos*-elements and *n* *Neg*-elements. The frequency of positives and negative classes is:

$$-P = p / (p + n)$$

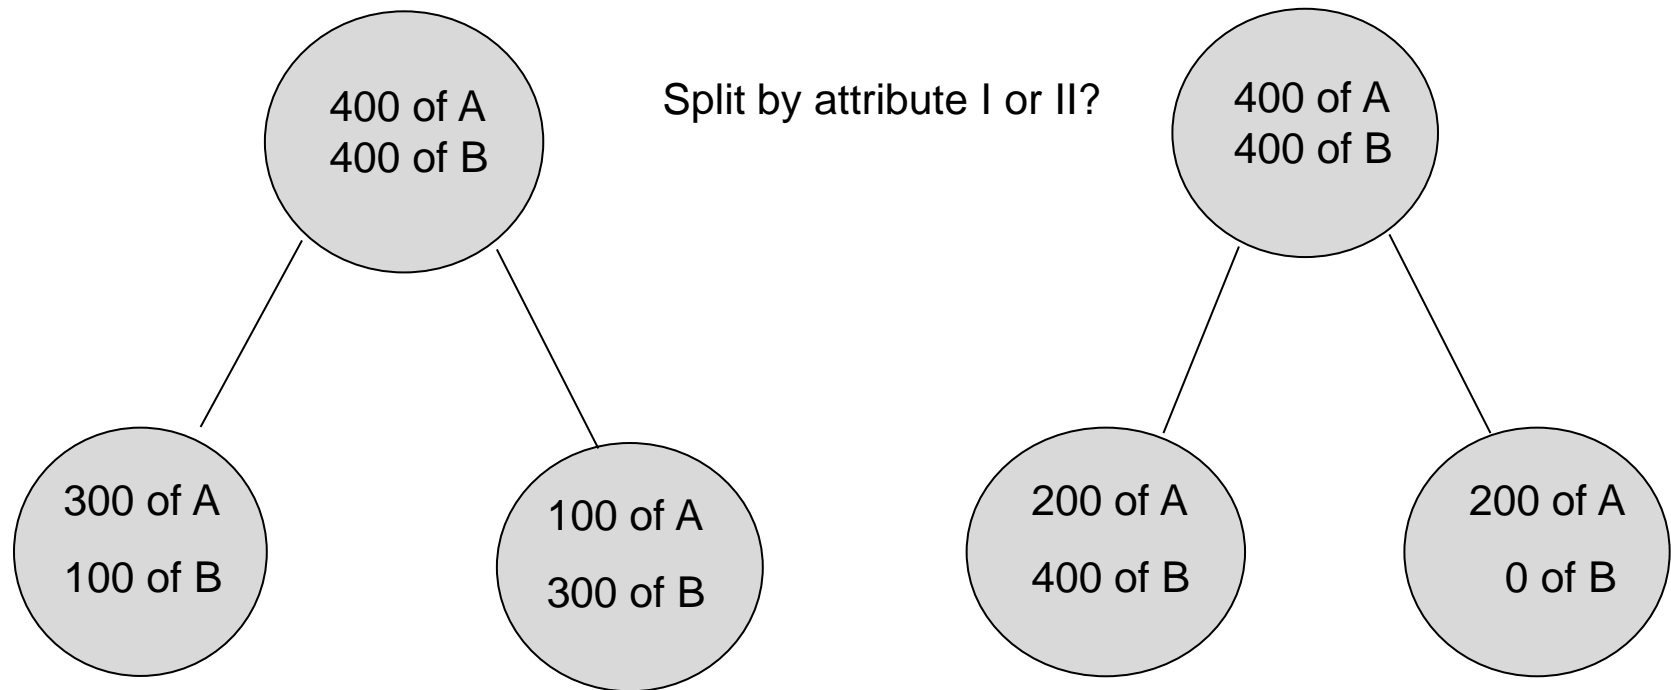
$$-N = n / (p + n)$$

$$Gini(S) = 1 - P^2 - N^2 \quad \in [0,0.5]$$

- If dataset *S* is split into *S*₁, *S*₂ then

$$Gini_{split}(S_1, S_2) = (p_1 + n_1)/(p + n) \cdot Gini(S_1) + \\ (p_2 + n_2)/(p + n) \cdot Gini(S_2)$$

Example



Example

$$\text{Gini}(p) = 1 - \sum_j p_j^2$$

Numbers of Cases		Proportion of Cases				Gini Index
A	B	A	B			
		p_A	p_B	p_A^2	p_B^2	$1 - p_A^2 - p_B^2$
400	400	0.5	0.5	0.25	0.25	0.5

Select the split that decreases the Gini Index most. This is done over all possible places for a split and all possible variables to split.

Gini Index Example

Number of Cases		Proportion of Cases				Gini Index	Info required	
A	B	A	B					
		p_A	p_B	p_A^2	p_B^2	$1 - p_A^2 - p_B^2$		
300	100	0.75	0.25	0.5625	0.0625	0.375	0.1875	0.5 * Gini(i)
100	300	0.25	0.75	0.0625	0.5625	0.375	0.1875	
						Total	0.375	
200	400	0.33	0.67	0.1111	0.4444	0.4444	0.3333	0.75 * Gini(i)
200	0	1	0	1	0	0	0	
						Total	<u>0.3333</u>	

Generate_DT(samples, attribute_list)

- Create a new node **N**;
- If **samples** are all of class **C** then label **N** with **C** and exit;
- If **attribute_list** is empty then label **N** with **majority_class(samples)** and exit;
- Select **best_split** from **attribute_list**;
- For each value **v** of attribute **best_split**:
 - Let **S_v** = set of samples with **best_split=v** ;
 - Let **N_v** = **Generate_DT**(**S_v**, **attribute_list \ best_split**) ;
 - Create a branch from **N** to **N_v** labeled with the test **best_split=v**;

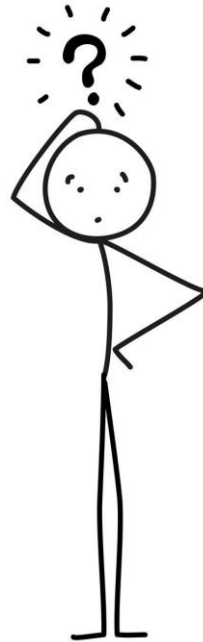
Time Complexity of Basic Algorithm

- Let m be the number of attributes
- Let n be the number of instances
- Assumption: Depth of tree is $\mathcal{O}(\log n)$
 - Usual assumption if the tree is not degenerate
- For each level of the tree all n instances are considered (best = v_i)
 - $\mathcal{O}(n \log n)$ work for a single attribute over the entire tree
- Total cost is $\mathcal{O}(mn \log n)$ since all attributes are eventually considered.
 - Without pruning (see next class)

Scalability of DT Algorithms

- Need to design for large amounts of data
- Two things to worry about:
 - Large number of attributes
 - Leads to a large tree
 - Takes a long time
 - Large amounts of data
 - Can the data be kept in memory?
 - Some new algorithms do not require all the data to be memory resident

How would you deal with missing values in the training data?



Outline for today

- Choosing a splitting attribute in decision trees
 - Information gain
 - Gain ratio
 - Gini index
- **Numeric attributes**
- Missing values
- Relational rules



C4.5 History

- The above procedure is the basis for Ross Quinlain's ID3 algorithm (so far works only for nominal attributes)
 - ID3, CHAID – 1960s
- The algorithm was improved and is now most widely used as C4.5 or C5.0 respectively, available in most DM software packages
 - Commercial successor: C5.0
- Witten et al. write “a landmark decision tree program that is probably the machine learning workhorse most widely used in practice to date”

C4.5 An Industrial-Strength Algorithm

- For an algorithm to be useful in a wide range of real-world applications it must:
 - Permit numeric attributes
 - Allow missing values
 - Be robust in the presence of noise
- Basic algorithm needs to be extended to fulfill these requirements

Numeric Attributes

- Unlike nominal attributes, every attribute has many possible split points
 - Standard method: binary splits
 - E.g. $\text{temp} < 45$
- Solution is straightforward extension:
 - Evaluate info gain (or other measure) for every possible split point of attribute
 - Choose “best” split point
 - Info gain for best split point is highest info gain for attribute
- Numerical attributes can be used several times in a decision tree, nominal attributes only once

Example

- Split on temperature attribute:

64	65	68	69	70	71	72	72	75	75	80	81	83	85
Yes	No	Yes	Yes	Yes	No	No	Yes	Yes	Yes	No	Yes	Yes	No

- E.g. temperature < 71.5 : yes/4, no/2
 temperature ≥ 71.5 : yes/5, no/3
- $\text{Info}([4,2],[5,3])$
 $= 6/14 \text{ info}([4,2]) + 8/14 \text{ info}([5,3])$
 $= 0.939 \text{ bits}$
- Place split points halfway between values

Binary Splits on Numeric Attributes

- Splitting (multi-way) on a nominal attribute exhausts all information in that attribute
 - Nominal attribute is tested (at most) once on any path in the tree
- Not so for binary splits on numeric attributes!
 - Numeric attributes may be tested several times along a path in the tree
- Disadvantage: tree is hard to read
- Remedy:
 - pre-discretize numeric attributes, or
 - allow for multi-way splits instead of binary ones using the Information Gain criterion

Outline for today

- Choosing a splitting attribute in decision trees
 - Information gain
 - Gain ratio
 - Gini index
- Numeric attributes
- **Missing values**
- Relational rules




Handling Missing Values / Training

- Ignore instances with missing values
 - Pretty harsh, and missing value might not be important
- Ignore attributes with missing values
 - Again, may not be feasible
- Treat missing value as another nominal value
 - Fine if missing a value has a significant meaning
- Estimate missing values
 - Data imputation: regression, nearest neighbor, mean, mode, etc.

Handling Missing Values / Classification


- Follow the leader
 - An instance with a missing value for a tested attribute (temp) is sent down the branch with the most instances



Instance included on the left branch

Handling Missing Values / Classification

- “Partition” the instance
 - branches show # of instances
 - Send down parts of the instance (e.g. 3/8 on Windy and 5/8 on Sunny) proportional to the number of training instances
 - Resulting leaf nodes get weighted in the result



Overfitting

- Two sources of abnormalities
 - Noise (randomness)
 - Outliers (measurement errors)
- Chasing every abnormality causes overfitting
 - Decision tree gets too large and complex
 - Good accuracy on training set, poor accuracy on test set
 - Does not generalize to new data any more
- Solution: prune the tree

Decision Trees - Summary


- Decision trees are a classification technique
- The output of decision trees can be used for descriptive as well as predictive purposes
- They can represent any function in the form of propositional logic
- Heuristics such as information gain are used to select relevant attributes
- Pruning is used to avoid overfitting

Outline for today

- Choosing a splitting attribute in decision trees
 - Information gain
 - Gain ratio
 - Gini index
- Numeric attributes
- Missing values
- **Relational rules**



The Shapes Problem



Shaded=standing
Unshaded=lying

Instances

Width	Height	Sides	Class
2	4	4	Standing
3	6	4	Standing
4	3	4	Lying
7	8	3	Standing
7	6	3	Lying
2	9	4	Standing
9	1	4	Lying
10	2	3	Lying

If width ≥ 3.5 and height < 7.0 then lying
If height ≥ 3.5 then standing

Classification Rules

If width ≥ 3.5 and height < 7.0 then lying

If height ≥ 3.5 then standing

- Work well to classify these instances ... but not necessarily for new ones.
- Problems?

Relational Rules

If width > height then lying

If height > width then standing

- Rules comparing attributes to constants are called propositional rules (propositional data mining)
- Relational rules are more expressive in some cases
 - define relations between attributes (relational data mining)
 - most DM techniques do not consider relational rules
- As a workaround for some cases, one can introduce additional attributes, describing if width > height
 - allows using conventional “propositional” learners

Propositional Logic

- Essentially, decision trees can represent any function in propositional logic
 - A, B, C: propositional variables
 - and, or, not, \Rightarrow (implies), \Leftrightarrow (equivalent): connectives
- A proposition is a statement that is either true or false
 - The sky is blue: color of sky = blue
- Decision trees are an example of a propositional learner.